

EECS 388



Introduction to Computer Security

Lecture 4:

Confidentiality

September 5, 2024

Prof. Halderman



Review: Message Integrity




Problem: Integrity of message from Alice to Bob over an *untrusted channel*

Approach: Alice must append bits to message that only Alice (or Bob) can make

Ideal solution: Random functions

Practical solution:



Pseudorandom functions (PRFs)

$f_k()$ is indistinguishable in practice from random, unless you know k

The **HMAC construction** turns a **cryptographic hash function** (e.g., **SHA-256**) into a **Message authentication code (MAC)**


For most practical purposes, we believe we can use **HMAC-SHA-256** as a PRF.

New Goal: Confidentiality



Confidentiality: Keep the content of message p secret from an *eavesdropper*

Approach and threat model:



Terminology:

p plaintext
 c ciphertext
 k secret key

$E_k()$ encryption function

$D_k()$ decryption function

Passive Eavesdropper

(a kind of **passive attacker**)

sees everything on the channel
but can't change anything

Digression: Historical Cryptography



Caesar cipher

First recorded use: Julius Caesar (100-44 BC)

Replace each plaintext *letter* with one a fixed number of places down the alphabet:

Encryption: $c_i := (p_i + k) \bmod 26$

Decryption: $p_i := (c_i - k) \bmod 26$

Examples using $k=3$: (that's the key Caesar used!)

p: ABCDEFGHIJKLMNOPQRSTUVWXYZ

+k: 33333333333333333333333333333333

=c: DEFGHIJKLMNOPQRSTUVWXYZABC

p: fox go wolverines

+k: 333 33 3333333333

=c: ira jr zroyhulqhv

[How would you break the Caesar cipher?]

Cryptanalysis of the Caesar cipher


Only 26 possible keys:

Try every possible k by “brute force”

Practical to do by hand!

How can a computer recognize the right one?

One solution: **Frequency analysis**. English text has characteristic letter frequency distribution:



Recognize with, e.g., chi-squared (χ^2) test

Digression: Historical Cryptography



Later advance: **Vigenère cipher** c. 1553
«*le chiffre indéchiffrable*» (“the indecipherable cipher”)

Encrypts successive letters using a sequence of Caesar ciphers keyed by the letters of a keyword.


For an n -letter “key word” k ,

Encryption: $c_i := (p_i + k_{i \bmod n}) \bmod 26$

Decryption: $p_i := (c_i - k_{i \bmod n}) \bmod 26$

Example: $k=ABC$ (i.e., $k_0=0$, $k_1=1$, $k_2=2$)

p :	bbbbb	amazon
$+k$:	<u>012012</u>	<u>012012</u>
$=c$:	bcdbcd	anczpp



[Break *le chiffre indéchiffrable*?]

Cryptanalysis of the Vigenère cipher

Easy, if we know the keyword length, n :

1. Divide ciphertext into n slices
2. Solve each slice as a Caesar cipher

How to find n ? One way: **Kasiski method**

Published 1863 by Kasiski (earlier known to Babbage?)

Repeated strings in long plaintext will sometimes, by coincidence, be encrypted with same key letters:

p :	CRYPTOISSHORTFORCRYPTOGRAPHY
$+k$:	<u>ABCDABCDABCDABCDABCDABCD</u>
$=c$:	CSASTPKVSIQUTGQCSASTPIUAQJB

Distance between repeats is (likely) a multiple of key size.
(ex: $16 \Rightarrow 16, 8, 4, 2, 1$). Use multiple repeats to narrow down

Another way: Iterate over n to find best match
(e.g., minimize sum of χ^2 for the individual Caesar ciphers)

One-time Pad (OTP)



Back to the present...

How can we achieve confidentiality securely?

One-time pad (OTP)

Alice and Bob share a secret,
very long string of random bits
(a “one-time pad”) k

Encryption: $c_i := p_i \oplus k_i$

Decryption: $p_i := c_i \oplus k_i$

XOR Facts

a	b	$a \oplus b$
0	0	0
0	1	1
1	0	1
1	1	0
$a \oplus b \oplus b = a$		
$a \oplus b \oplus a = b$		

Pro: Provably secure

Information-theoretically secure

(no computational complexity assumptions)

First proof published by Claude Shannon in 1949


[Show Mallory can't do better than guessing]



Con: Usually impractical [Why?] [Exceptions?]

Caution “one-time” means you MUST NEVER
reuse any part of the pad

If you do: Let k_i be pad bit. From ciphertexts $(a \oplus k_i)$
and $(b \oplus k_i)$, attacker learns $a \oplus b$. [How's that useful?]



Stream Cipher: Use a PRG for Confidentiality **M**

More practical approach to confidentiality:

Use a **pseudorandom generator (PRG)** instead of a truly random pad


Recall: A PRG $g_k()$ is practically indistinguishable from a random stream of bits, unless you know k .

Called a **stream cipher**

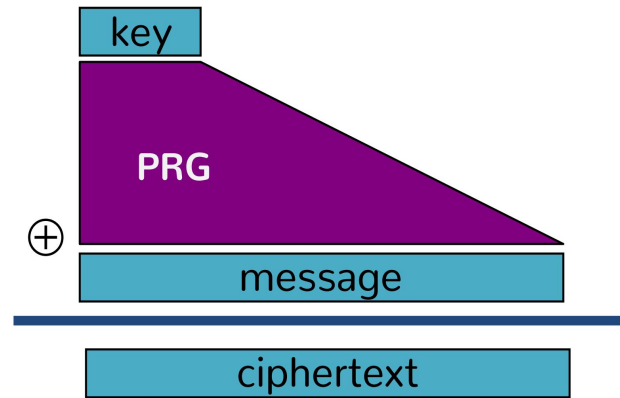
Alice and Bob choose PRG $g()$, share secret key k

Encryption: $s := g_k()$; $c_i := p_i \oplus s_i$

Decryption: $s := g_k()$; $p_i := c_i \oplus s_i$



Caution! NEVER reuse keys
NEVER reuse PRG output bits



Provably secure if $g()$ is a secure PRG, under complexity assumptions

However... we don't know how to prove that secure PRGs even exist!

Best we can do (again): Use well studied functions where we haven't spotted a problem yet

Examples: **RC4** **ChaCha20**

Block Ciphers



Another approach:

Block ciphers

consist of a function that **encrypts** fixed-size (n -bit) blocks with a reusable key k :

$$E_k(p) : \{0,1\}^{|k|} \times \{0,1\}^n \rightarrow \{0,1\}^n$$

and an inverse function that **decrypts** the blocks when used with same key:

$$D_k(c) = E_k^{-1}(c) : \{0,1\}^{|k|} \times \{0,1\}^n \rightarrow \{0,1\}^n$$

such that $\forall k : D_k(E_k(p)) = p$.

In effect, k selects one *permutation* from the set of $2^n!$ possible permutations of E 's domain.

A block cipher is *different* from a PRF. [Why?]

What do we want, if not a PRF?

Pseudorandom permutation (PRP)

A secure PRP is a function that cannot practically be distinguished from a truly random permutation unless you know k . (Similar to the PRF game)

Annoying question again:

Do PRPs *actually* exist?

Same annoying answer:

We don't know. :(

Best we can do:

Design a complex function that's invertible if and (hopefully) only if you know k

Examples: ~~DES~~ AES

AES Block Cipher



Today's most common block cipher:

AES (Advanced Encryption Standard)

aka **Rijndael**, for its designers, Rijmen and Daemen

- Standardized by **NIST** in 2001 after winning a long, public international design competition
- Efficient in both software and hardware.
Hardware-accelerated in many modern CPUs
- Widely believed to be a secure PRP
(but we don't know how to prove it)

Fixed block size: 128 bits

Variable key size: 128, 192, or 256 bits

10, 12, or 14 **rounds** (based on key size)

Generates $r := \text{\#rounds subkeys}$ from k .

Performs same set of operations r times,
each with a different subkey

Each AES round

128-bits input

128-bit subkey

128-bit output

picture as operations on
a 4×4 grid of 8-bit values

$S_{0,0}$	$S_{0,1}$	$S_{0,2}$	$S_{0,3}$
$S_{1,0}$	$S_{1,1}$	$S_{1,2}$	$S_{1,3}$
$S_{2,0}$	$S_{2,1}$	$S_{2,2}$	$S_{2,3}$
$S_{3,0}$	$S_{3,1}$	$S_{3,2}$	$S_{3,3}$

Four steps:

1. Non-linear substitution

Run each byte thru a nonlinear function (lookup table)

2. Shift rows

Circular-shift each row: i^{th} row shifted by i (0–3)

3. Linear-mix columns

Treat each column as a 4-vector;
multiply by a constant invertible matrix

4. Key-addition

XOR each byte with corresponding byte of round subkey

To decrypt, just undo the steps, in reverse order

Padding and Block Cipher Modes



Challenge for block ciphers:

How to encrypt **arbitrary-sized messages**?

Padding: Add bytes to end of message to make it a multiple of block size

Flawed approach: add zeros [What's the issue?]

| MM MM MM MM MM 00 00 00 |

Don't know what to remove after decryption!

Better approach (**PKCS7**): Add **n** bytes of value **n**

| MM MM MM MM MM 03 03 03 |

Edge case: Message that ends at block boundary?

| MM MM MM MM MM MM MM MM | 08 08 08 08 08 08 08 08 |

Add an **entire block** of padding

Ensures receiver can **unambiguously** distinguish the padding from the message after decrypting

Cipher modes: Algorithms for applying block ciphers to more than one block

Flawed approach: [What's the issue?]

Encrypted codebook (ECB) mode

Simply encrypt each block independently: $c_i := E_k(p_i)$



Plaintext

Pseudorandom

ECB mode

More Cipher Modes



Cipher-block chaining (CBC) mode

“Chains” ciphertexts to obscure later ones

Choose a **random initialization vector IV**

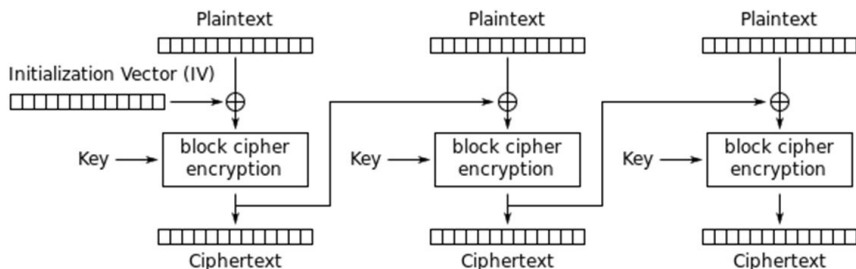
Encrypt: $c_0 := IV; c_i := E_k(p_i \oplus c_{i-1})$

Decrypt: $p_i := D_k(c_i) \oplus c_{i-1}$,

[Why do we need the IV?]

Have to send IV with ciphertext

Can't encrypt blocks in parallel or out of order



Counter (CTR) mode

Turns a block cipher into a stream cipher

Generate **keystream** s for k and **unique nonce**:


$s := E_k(\text{nonce}||0) \parallel E_k(\text{nonce}||1) \parallel E_k(\text{nonce}||2) \parallel \dots$

Encrypt: $c := p \oplus s$ Decrypt: $p := c \oplus s$

Benefits: Doesn't require padding

Efficient parallelism/random access

Caution: Never reuse **nonce** for same **k**!



Getting both confidentiality and integrity?



Integrity (**tampering**)

Let $f()$ be a secure PRF.

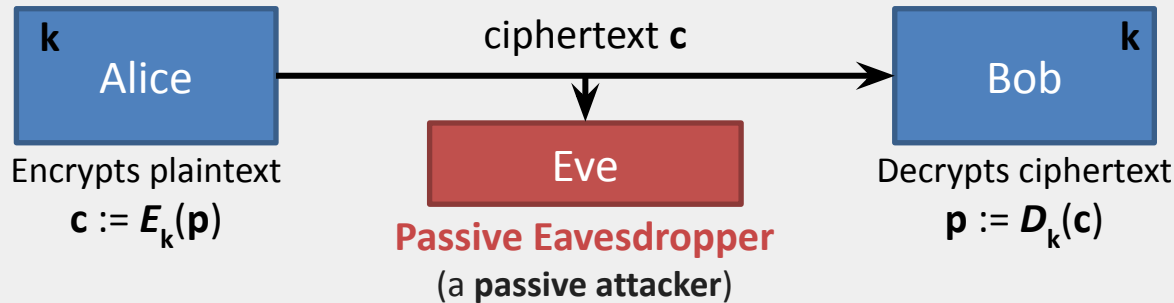
In practice: e.g., **HMAC-SHA-256**



Confidentiality (**eavesdropping**)

Construct $E()$ and $D()$ from secure PRG (a stream cipher) **or** secure PRP (a block cipher) with appropriate padding/cipher mode.

In practice: e.g., **AES-128 in CTR mode**



What if we want integrity and confidentiality *at the same time*?
(Next lecture!)

Coming Up



Reminders:

Lab Assignment 1 due TODAY at 6 p.m.

Quiz on Canvas after every lecture

Project 1, Part 1 due next Thursday at 6 p.m.

Tuesday

Combining Confidentiality and Integrity

Confidentiality attacks,
authenticated encryption

Thursday

Public Key Cryptography

Diffie-Hellman key exchange,
RSA encryption,
digital signatures