

EECS 388



# Introduction to Computer Security

Lecture 16:

## Control Hijacking (Part 2)

October 24, 2024  
Aidan Delwiche

*champ gives the craziest side eye...*



# example.s (x64)




```
void foo(char *str) {  
    char buffer[4];  
    strcpy(buffer, str);  
}
```

foo:

```
    push    rbp  
    mov     rbp, rsp  
    sub     rsp, 4  
    mov     rsi, rdi  
    lea     rdi, [rbp-4]  
    call    strcpy  
    leave  
    ret
```

str\_ptr: "AAAABBBBBBBB1234567"



# example.s (x64)



```
void foo(char *str) {  
    char buffer[4];  
    strcpy(buffer, str);  
}
```


foo:

```
    push    rbp  
    mov     rbp, rsp  
    sub     rsp, 4  
    mov     rsi, rdi  
    lea     rdi, [rbp-4]  
    call    strcpy  
    leave  
    ret
```

call strcpy



str\_ptr: "AAAABBBBBBBB1234567"



# Stack Shellcode



Good

Evil

DEADLY

*buffer*

*“AAAA”*

*0xEBFE4141...*

0xffffffff88888888

0xffffffff8888

000111b1 55

000111b3 48 89 e5

000111b6 eb fe

undefined `main()`

PUSH RBP

MOV RBP, RSP

LAB\_000111b4

JMP LAB\_000111b4


*0x41414141...*

*0xFFFFFFFF88888888*

*prev FP*


*prev FP*

*prev FP*



## Write XOR Execute

- Write (heap/stack)
- Execute (code segments)
- Never both!



```
int delete_account(char* username,  
    int length, VOID* creds) {  
    int admin;  
    char name[100];  
    admin = check_admin(creds);  
    strncpy(name, username, length);  
    canonicalize_username(name);  
    if (admin) {delete_user(name);}  
    return (admin > 0);  
}
```



# Cat-and-Mouse Exploitation

DEP

Buffer Overflow  
Stack Shellcode

Data-only attacks  
**Return-to-libc**

Defender:

DEP prevents executing injected shellcode

Attacker:

Reuse code that already exists

Already marked as executable!





# Recall... Calling Convention



## Caller:

Push arguments to registers **RDI**, **RSI**, **RDX**, **RCX**, **R8**, **R9**

Push **rip**

Jump

## Callee:

Push **rbp**



Move **rbp** to **rsp**

# Stack after entering a new function



**foo:**

```
push    rbp
mov     rbp, rsp
sub     rsp, 0x10
leave
ret
```

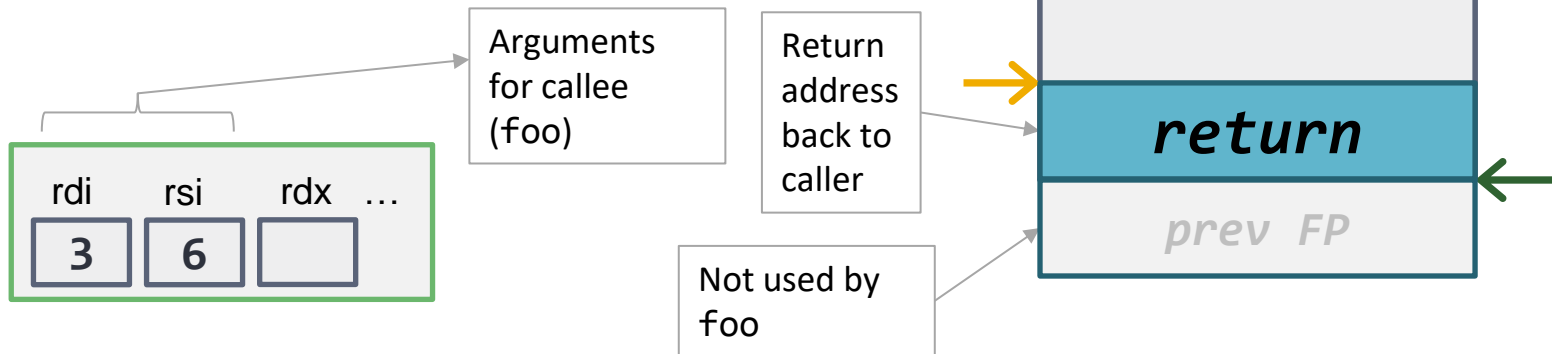


# Stack after entering a new function



**foo:**

```
push    rbp
mov     rbp, rsp
sub     rsp, 0x10
leave
ret
```



# Stack after entering a new function



(some function that already exists)


## ATTACK:

```
push rbp
mov rbp, rsp
sub rsp, 0x10
leave
ret
```

Arguments  
for callee  
(ATTACK)

Return  
address  
back to  
caller

Not used by  
ATTACK



# Stack after entering a new function



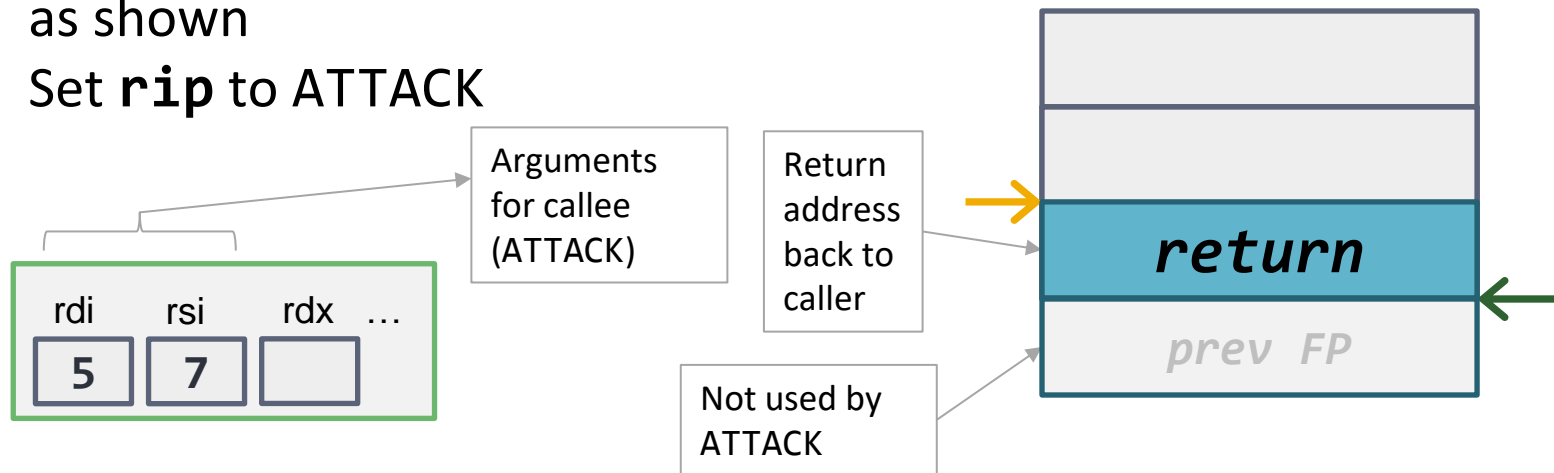
Our goal:

Execute ATTACK using a `ret` instead of a `call`

“Return-to-ATTACK”

How?

1. Set stack, **rsp**, and argument registers as shown
2. Set **rip** to ATTACK



# Stack after entering a new function



Our goal:

Execute ATTACK using a `ret` instead of a `call`

“Return-to-ATTACK”

How?

0. Enter vulnerable function
1. Set stack, **rsp**, and argument registers as shown
2. Set **rip** to ATTACK

```
void foo(int a, int b) {  
    char buf1[16];  
}
```

```
void vulnerable() {  
    int i = 3;  
    int j = 6;  
    char buffer[8];  
    gets(buffer);  
    foo(i, j);  
}
```

vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call    foo
leave
ret
```

# Attack surface



vulnerable:

```
push    rbp
mov     rbp, rsp
```

```
push    3
```

```
push    6
```

```
sub     rsp, 8
```

```
lea     rdi, [rbp - 16]
```

```
call    gets
```


```
mov     rsi, [rbp - 8]
```

```
mov     rdi, [rbp - 4]
```

```
call    foo
```

```
leave
```

```
ret
```






# Attack surface



vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call   gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call    foo
leave
ret
```




# Attack surface



vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call   foo
leave
ret
```





# Attack surface



vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call    foo
leave
```

**ret**



# Attack surface



**vulnerable:**


```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call    foo
leave
```

**ret**

pop **rip**



After executing the ret instruction, this part of the stack will be gone!



# Rewind to before gets call



Our goal

Return address  
back to caller

Arguments  
for callee



*return*



*prev FP*

rdi rsi rdx ...

5

7

rdi rsi rdx ...

buf\_ptr

After executing the ret instruction, this part of the stack will be gone!



*buffer*

6

3

*main FP*

*return*

*prev FP*



# Rewind to before foo call



Our goal

Return address  
back to caller

Arguments  
for callee



*return*



*prev FP*

rdi rsi rdx ...

5

7

rdi

rsi

rdx ...

3

6

After executing the `ret` instruction, this part of the stack will be gone!



*buffer*

6

3

*main FP*

*return*

*prev FP*

# On your marks, get set, ...






After executing the ret instruction, this part of the stack will be gone!

Our goal


Return address  
back to caller

Arguments  
for callee

Careful! What should  
these two 'return's be?




# Attacker 2.0 – during gets call

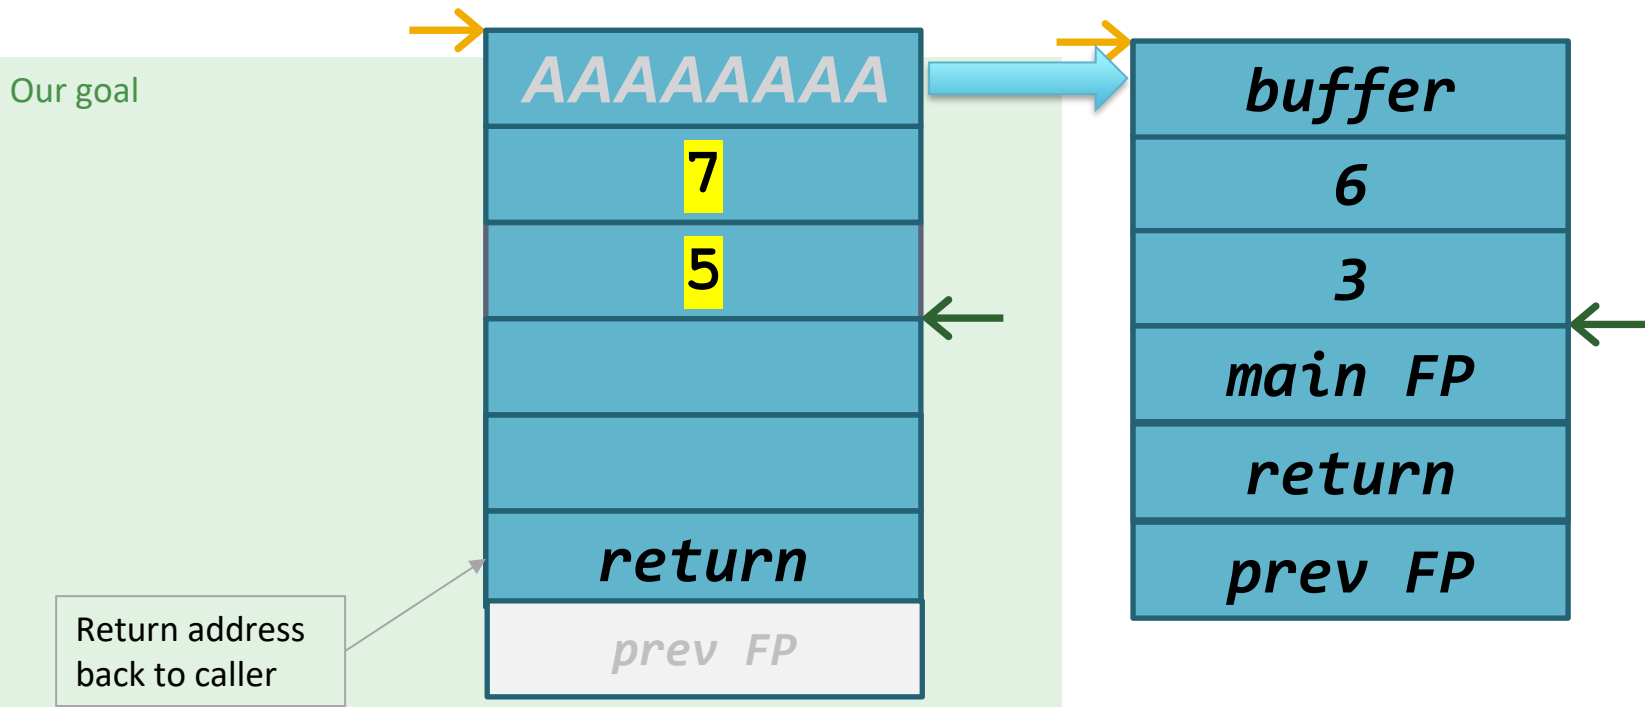




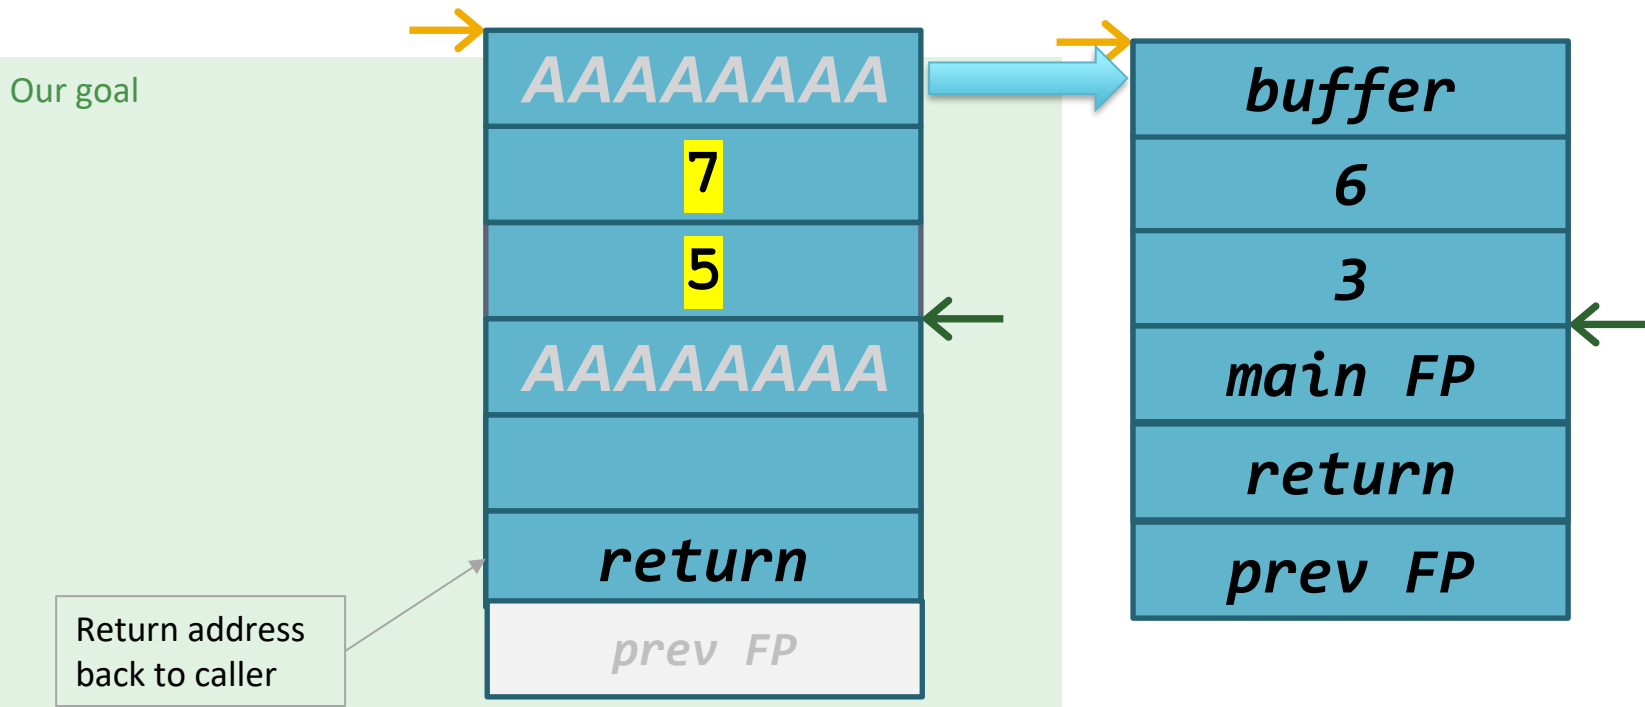
# Attacker 2.0 – during gets call



# Attacker 2.0 – during gets call



# Attacker 2.0 – during gets call



# Attacker 2.0 – during gets call



0x000000ffa804fef0

ATTACK:

```
push rbp  
mov rbp, rsp
```

```
...  
leave  
ret
```

Our goal

AAAAAAAA

7

5

AAAAAAAA

0x000000ffa804fef0

*return*

*prev FP*

Return address  
back to caller

*buffer*

6

3

*main FP*

*return*

*prev FP*

# Attacker 2.0 – during gets call



0x000000ffa804fef0

ATTACK:

```
push rbp  
mov rbp, rsp
```

```
...  
leave  
ret
```

Our goal

AAAAAAAA

7

5

AAAAAAAA

0x000000ffa804fef0

AAAAAA  
return

Return address  
back to caller

prev FP

buffer

6

3

main FP

return


prev FP

# Attack surface after gets



vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call    foo
leave
ret
```





# Attack surface after gets



vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call   foo
leave
ret
```



# Attack surface after gets




vulnerable:

```
push    rbp
mov     rbp, rsp
push    3
push    6
sub     rsp, 8
lea     rdi, [rbp - 16]
call    gets
mov     rsi, [rbp - 8]
mov     rdi, [rbp - 4]
call    foo
leave
```

**ret**

pop **rip**

rdi	rsi	rdx	...
5	7		



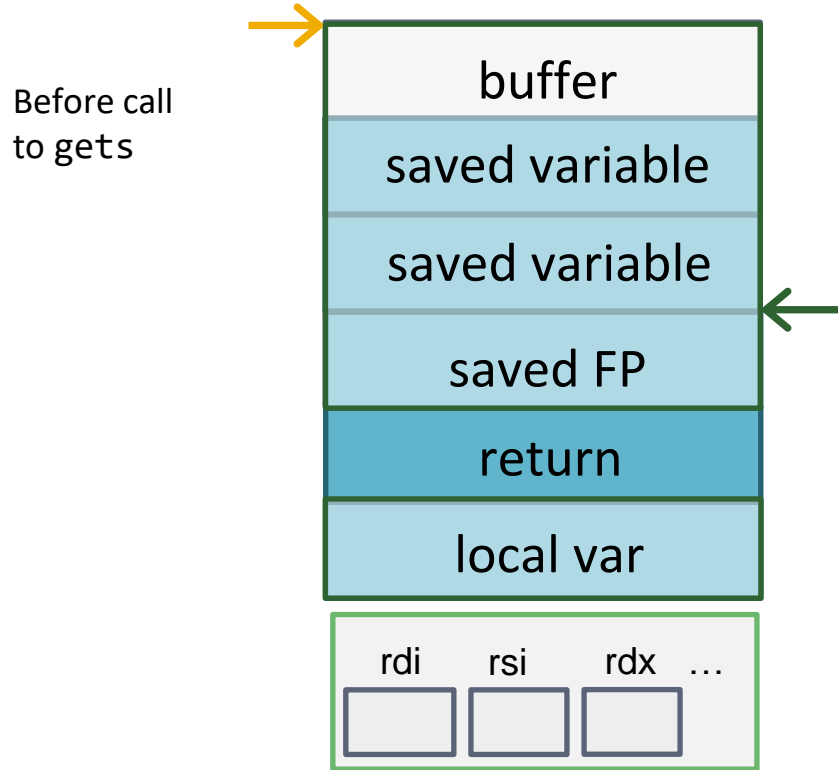
← 0x4141414141414141




# Let's see it in action



Return to libc attack setup




Normal function call setup




# Let's see it in action



Return to libc attack setup




Normal function call setup



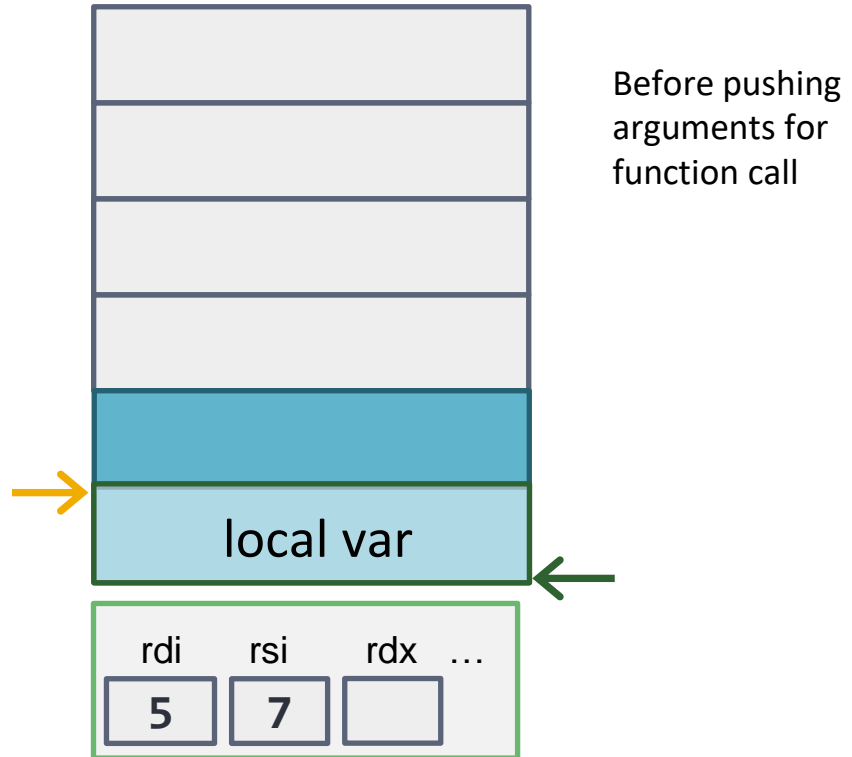
# Let's see it in action



Return to libc attack setup




Normal function call setup




# Let's see it in action



Return to libc attack setup




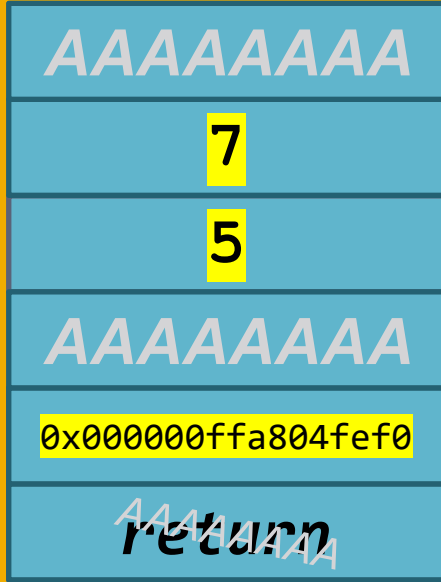
Normal function call setup





# Let's see it in action





Our exploit



Normal function call setup



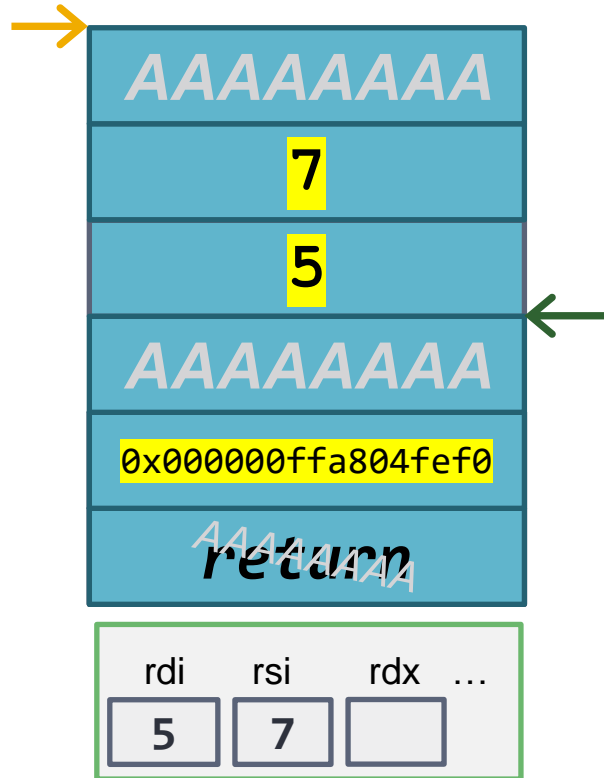
Normal function call setup



# Let's see it in action



Return to libc attack setup



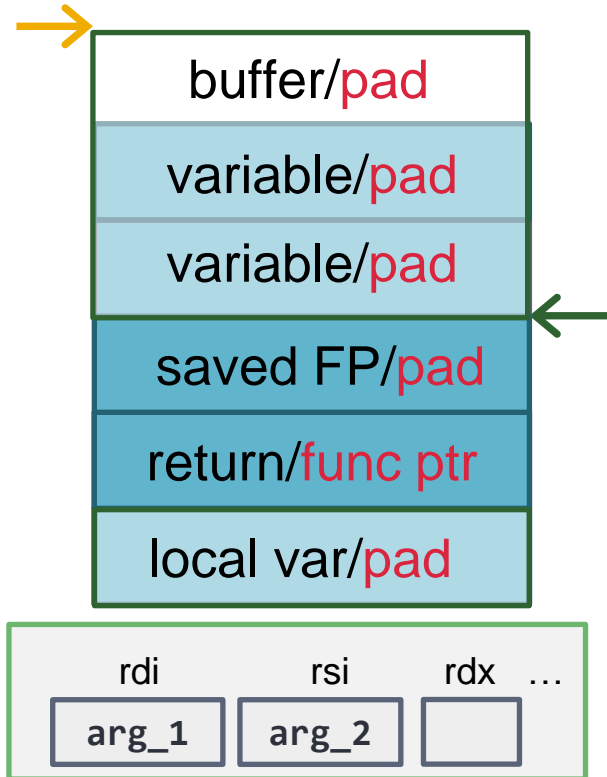
Normal function call setup



# Let's see it in action



Return to libc attack setup




Normal function call setup



# Let's see it in action



Return to libc attack setup



Normal function call setup


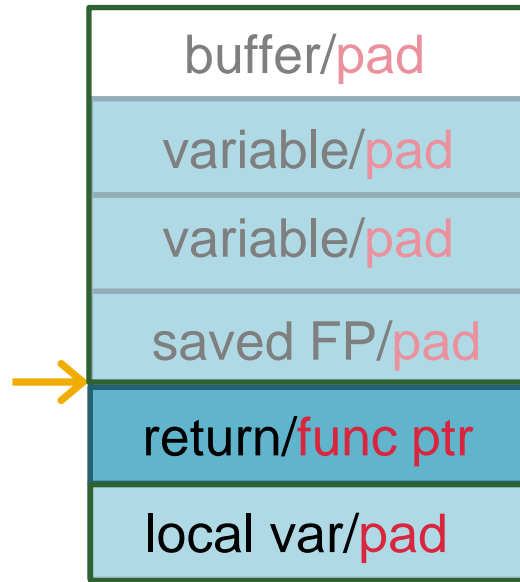




# Let's see it in action





Return to libc attack setup



**pad** ←

Normal function call setup




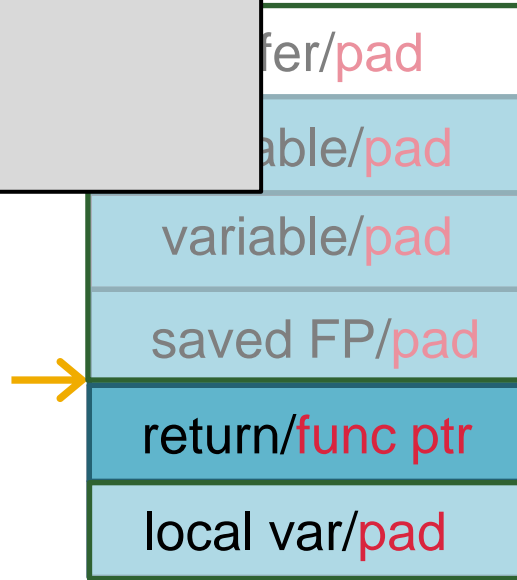
# Let's see it in action





## ATTACK:

```
push rbp
mov rbp, rsp
...
leave
ret
```

## libc attack setup



## Normal function call setup




# Let's see it in action




## ATTACK:

```
push rbp
mov rbp, rsp
...
leave
ret
```



## libc attack setup



pad ←



## Normal function call setup




# Let's see it in action




## ATTACK:

```
push rbp
mov rbp, rsp
...
leave
ret
```



## libc attack setup



pad ←



## Normal function call setup



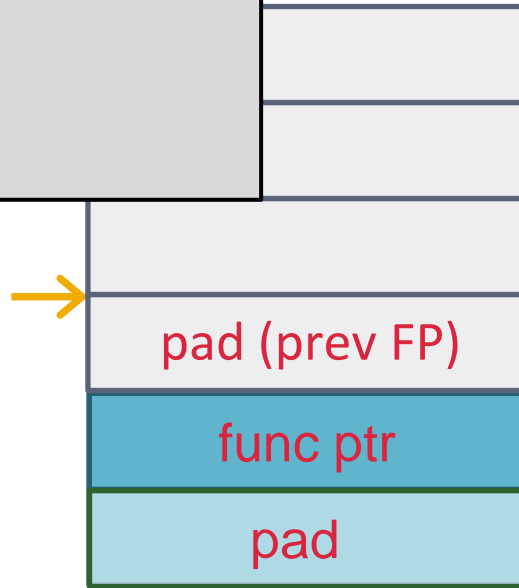
# Let's see it in action




ATTACK:

```
push rbp
mov rbp, rsp
...
leave
ret
```



libc attack setup



pad ←



Normal function call setup





# Let's see it in action





ATTACK:

```
push rbp
mov rbp, rsp
...
leave
ret
```

libc attack setup



Normal function call setup



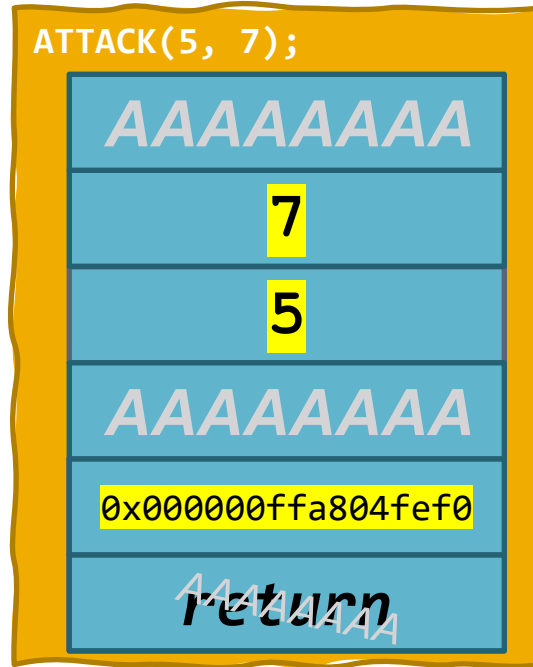

Invoke any function that exists in the binary  
execv is a popular one

The **execv()**, **execvp()**, and **execvpe()** functions provide an array of pointers to null-terminated strings that represent the argument list available to the new program. The first argument, by convention, should point to the filename associated with the file being executed. The array of pointers *must* be terminated by a NULL pointer.


# Return to libc



```
int main() {  
    char *args[] = {"/bin/ls", NULL};  
    execv("/bin/ls", args);  
}
```



Exercise: Walk through previous example using the below payload, and assuming foo takes double length values instead of int.





```
#include <sys/mman.h>
```

```
int mprotect(void *addr, size_t len, int prot);
```

## Description

The *mprotect()* function shall change the access protections to be that specified by *prot* for those whole pages containing any part of the address space of the process starting at address *addr* and continuing for *len* bytes. The parameter *prot* determines whether read, write, execute, or some combination of accesses are permitted to the data being mapped. The *prot* argument should be either PROT\_NONE or the bitwise-inclusive OR of one or more of PROT\_READ, PROT\_WRITE, and PROT\_EXEC.

# Return to libc: defense



Defender:

The attackers are calling other functions

Defender:

Let's remove functions we don't need!



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

DEP

**Extraneous function removal**



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

DEP

Extraneous function removal

**Return Oriented Programming (ROP)**

Defender:

Take out functions that can launch shells

Attacker:

Use the instructions that are still there

# Return Oriented Programming

Return to libc without calling full functions

Build arbitrary functionality via “gadgets”

Turing complete

## The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)

Hovav Shacham<sup>\*</sup>  
Department of Computer Science & Engineering  
University of California, San Diego  
La Jolla, California, USA  
hovav@hovav.net

### ABSTRACT

We present new techniques that allow a return-into-libc attack to be mounted on x86 executables that calls *no functions at all*. Our attack combines a large number of short instruction sequences to build *gadgets* that allow arbitrary computation. We show how to discover such instruction sequences by means of static analysis. We make use, in an essential way, of the properties of the x86 instruction set.

using the short sequences we find in a specific distribution of GNU libc, and we conjecture that, because of the properties of the x86 instruction set, in any sufficiently large body of x86 executable code there will feature sequences that allow the construction of similar gadgets. (This claim is our *thesis*.) Our paper makes three major contributions:

1. We describe an efficient algorithm for analyzing libc to recover the instruction sequences that can be used in

Small section of code

Contains a very small number of instructions

Ends in a **ret**

**Not an existing function body**

arg[10] = 0x00


var = var - 10

foo:

```
push rbp
mov rsp, rbp
mov rax, [rbp + 4]
add rax, 10
mov [rax], 0x00
sub rax, 10
leave
ret
```

foo+0x20:

```
sub rax, 10
leave
ret
```






**Variable-length instructions:** The achilles heel of x86.

<u>01 01</u>	<u>c3</u>
add [rcx], rax	ret

**Variable-length instructions:** The achilles heel of x86.



**Variable-length instructions:** The achilles heel of x86.



**Bytes in the Code Section:**

00 F7 C7 07 00 00 00 0f 95 53 c3

Full Gadget:

00 F7 C7 07 00 00 00 0f 95 53 c3

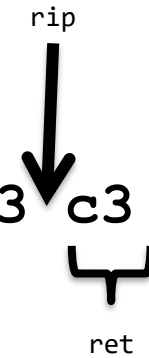




Diagram illustrating a Return Oriented Point (ROP) gadget. The sequence of bytes is 00 F7 C7 07 00 00 00 0f 95 53 c3. The instruction pointer (rip) points to the byte c3, which is the start of the instruction ret.

Full Gadget:  
ret




Full Gadget:  
push rbx  
ret



Full Gadget:

```
xchg rbp, rax
push rbx
ret
```



Full Gadget:  
 setne [rbx-61]  
 <no return>



rip  
↓  
00 F7 C7 07 00 00 00 0f 95 53 c3

Full Gadget:

<none - invalid instruction>

rip  
↓  
00 F7 C7 07 00 00 00 0f 95 53 c3

Full Gadget:

<none - invalid instruction>

rip  
↓  
00 F7 C7 07 00 00 00 0f 95 53 c3


Full Gadget:

<none - invalid instruction>

rip  
↓  
00 F7 C7 07 00 00 00 0f 95 53 c3

Full Gadget:


<none - invalid instruction>



Full Gadget:


```

mov rdi, 0x0F000000
xchg rbp, rax
push rbx
ret
  
```



Full Gadget:

```
test rdi, 0x00000007
setne [rbx-61]
<no return>
```



Full Gadget:

```
add bh, dh
mov rdi, 0xF000000
xchg rbp, rax
push rbx
ret
```

```

8052867: 74 1c          je      8052885 <__sprintf_chk@plt+0x8ab5>
8052869: 8b c4 04      mov     ebp,DWORD PTR [ebp+0x4]
805286c: 83 c3 01      add     ebx,0x1
805286f: 85 ed        test    ebp,ebp
8052871: 75 e5        jne     8052858 <__sprintf_chk@plt+0x8a88>
8052873: 8b 54 24 30    mov     edx,DWORD PTR [esp+0x30]
8052877: 83 44 24 0c 08 add     DWORD PTR [esp+0xc],0x8
805287c: 8b 44 24 0c    mov     eax,DWORD PTR [esp+0xc]
8052880: 39 42 04      cmp     DWORD PTR [edx+0x4],eax
8052883: 77 bf        ja      8052844 <__sprintf_chk@plt+0x8a74>
8052885: 83 c4 1c      add     esp,0x1c
8052888: 89 d8        mov     eax,ebx
805288a: 5b          pop     ebx
805288b: 5e          pop     esi
805288c: 5f          pop     edi
805288d: 5d          pop     ebp
805288e: c3          ret
805288f: 90          nop
8052890: 56          push    esi
8052891: 53          push    ebx
8052892: 31 d2        xor     edx,edx
8052894: 8b 5c 24 0c    mov     ebx,DWORD PTR [esp+0xc]
8052898: 8b 74 24 10    mov     esi,DWORD PTR [esp+0x10]
805289c: 0f b6 0b      movzx   ecx,BYTE PTR [ebx]
805289f: 84 c9        test    cl,cl
80528a1: 74 1c        je      80528bf <__sprintf_chk@plt+0x8aef>
80528a3: 90          nop
80528a4: 8d 74 26 00    lea     esi,[esi+eiz*1+0x0]
80528a8: 89 d0        mov     eax,edx
80528aa: 83 c3 01      add     ebx,0x1
80528ad: c1 e0 05      shl     eax,0x5
80528b0: 29 d0        sub     eax,edx
80528b2: 31 d2        xor     edx,edx
80528b4: 01 c8        add     eax,ecx

```

0xc3 : ret

Could be part of another instruction!

Empirically, about one in every 178 bytes



Gadget 1:

```
mov rax, 0x10; ret
```

Gadget 2:

```
add rax, rbp; ret
```

Gadget 3:

```
mov [rax+8], rax;  
ret
```

Gadget 4:

```
mov rbp, rsp; ret
```

# ROP chains



Gadget 1:

```
mov rax, 0x10; ret
```

Gadget 2:

```
add rax, rbp; ret
```

Gadget 3:

```
mov [rax+8], rax;  
ret
```

Gadget 4:

```
mov rbp, rsp; ret
```

<i>buffer</i>
<i>Local var</i>
<i>saved FP</i>
<i>ret</i>
<i>Local var</i>
<i>prev FP</i>
<i>prev ret</i>

# ROP chains



Gadget 1:

```
mov rax, 0x10; ret
```

Gadget 2:

```
add rax, rbp; ret
```

Gadget 3:

```
mov [rax+8], rax;  
ret
```

Gadget 4:

```
mov rbp, rsp; ret
```


<i>pad</i>
<i>pad</i>
<i>pad</i>
<i>&amp;gadget 1</i>
<i>&amp;gadget 2</i>
<i>&amp;gadget 3</i>
<i>&amp;gadget 4</i>

# ROP chains



ROP Chain:

```
mov rax, 0x10; ret
add rax, rbp; ret
add rax, rbp; ret
mov [rax+8], rax; ret
mov rbp, rsp
ret
```




# ROP chains



ROP Chain:

```
mov rax, 0x10; ret
add rax, rbp; ret
add rax, rbp; ret
mov [rax+8], rax; ret
mov rbp, rsp
ret
```




# ROP chains



ROP Chain:

```
mov rax, 0x10; ret
add rax, rbp; ret
add rax, rbp; ret
mov [rax+8], rax; ret
mov rbp, rsp
ret
```



# ROP chains



ROP Chain:

```
mov rax, 0x10; ret
add rax, rbp; ret
add rax, rbp; ret
mov [rax+8], rax; ret
mov rbp, rsp
ret
```




# ROP chains



ROP Chain:

```
mov rax, 0x10; ret
add rax, rbp; ret
add rax, rbp; ret
mov [rax+8], rax; ret
mov rbp, rsp
ret
```







# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

DEP

Extraneous function removal

**ASLR**

Defender:

We can't take out all the rets from our code

Defender:

Let's just move around where the code lives

Address Space Layout Randomization


Make it extremely hard to predict references

Requires many changes to compilation and/or loading

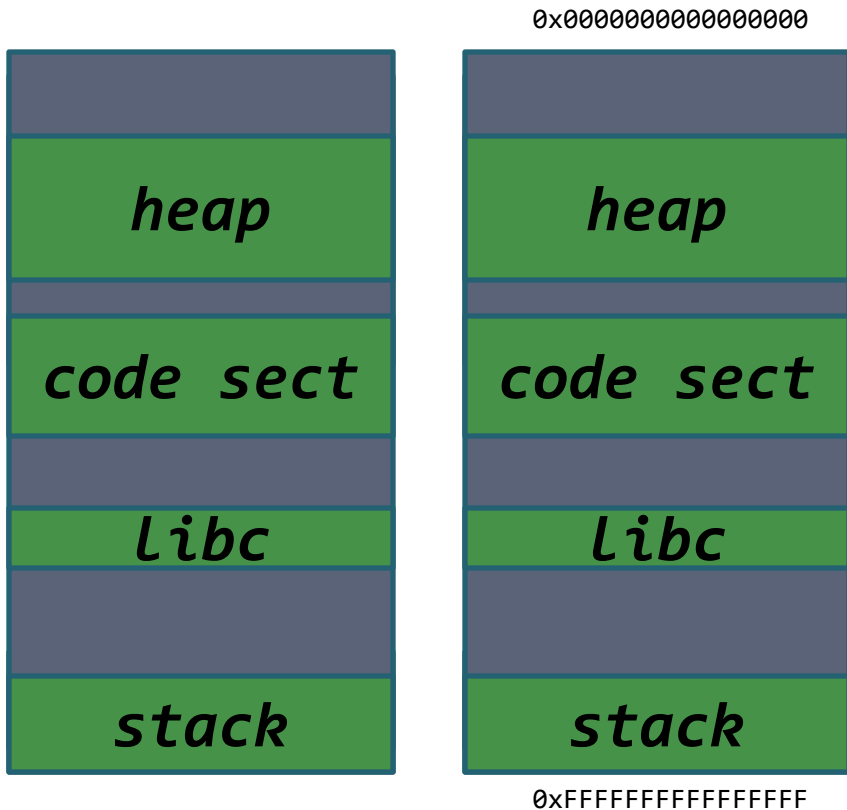
Code must be “relocatable” or “position independent”

<Details are out-of-scope>

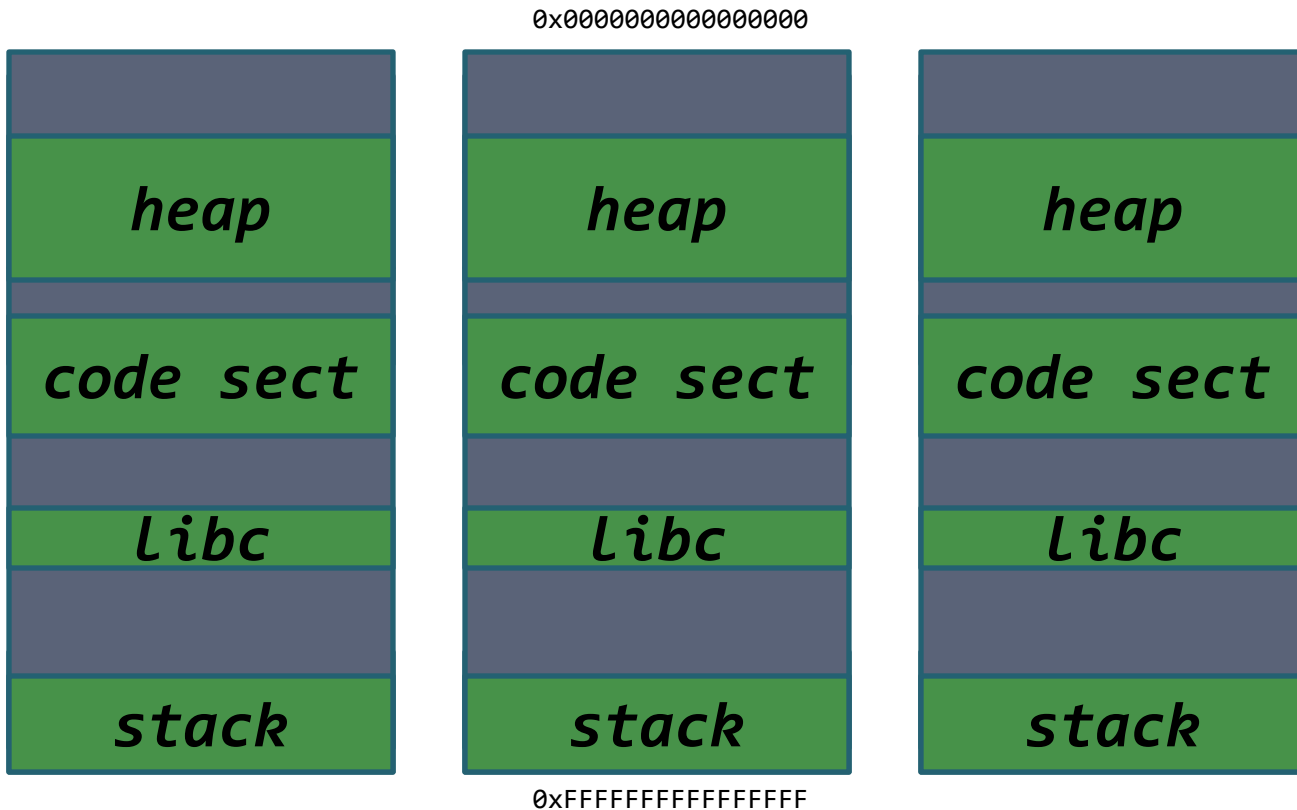
# Memory Layout (no ASLR)




# Memory Layout (no ASLR)



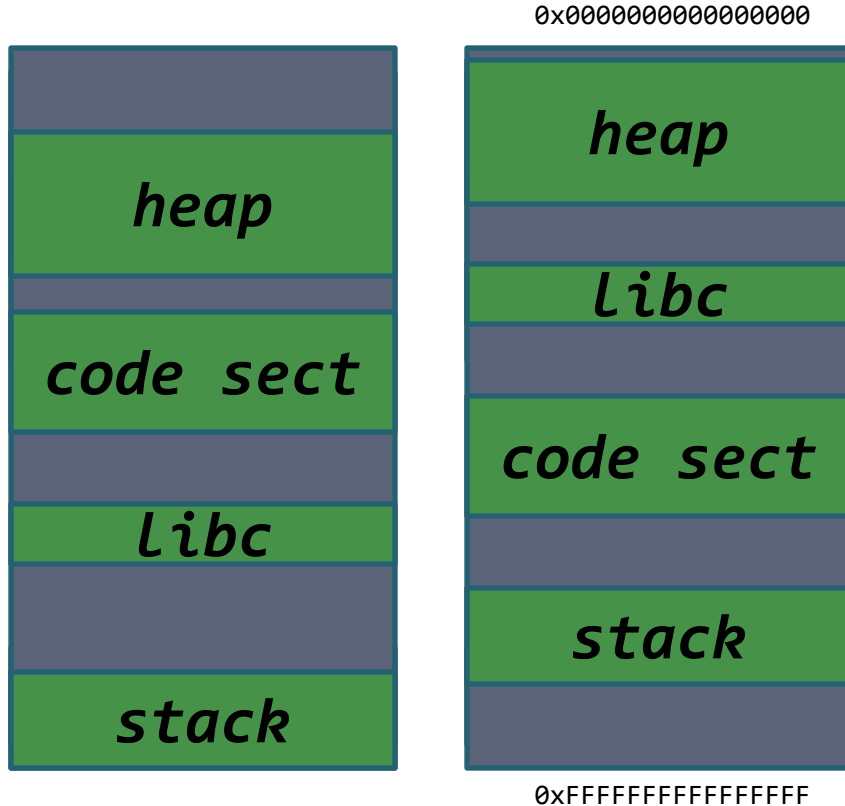
# Memory Layout (no ASLR)



# Memory Layout (with ASLR)

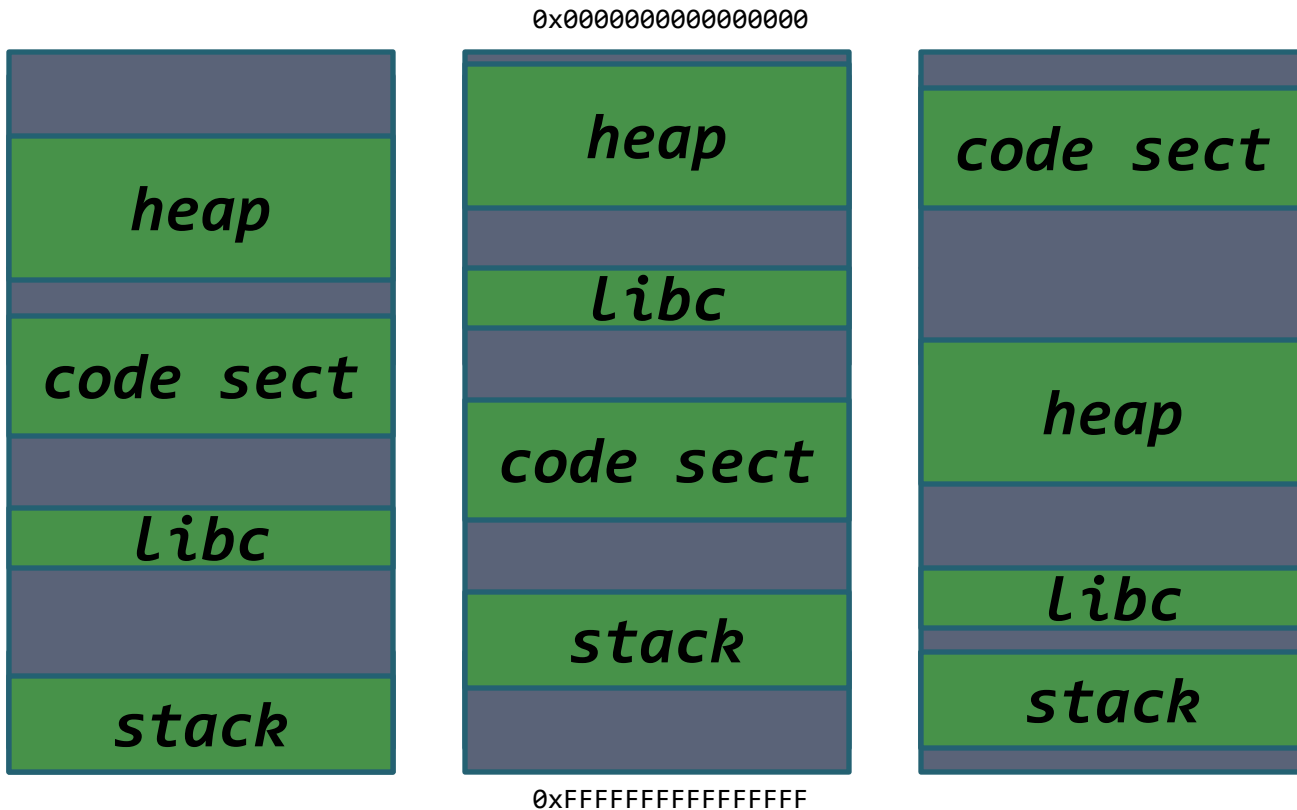


# Memory Layout (with ASLR)





# Memory Layout (with ASLR)



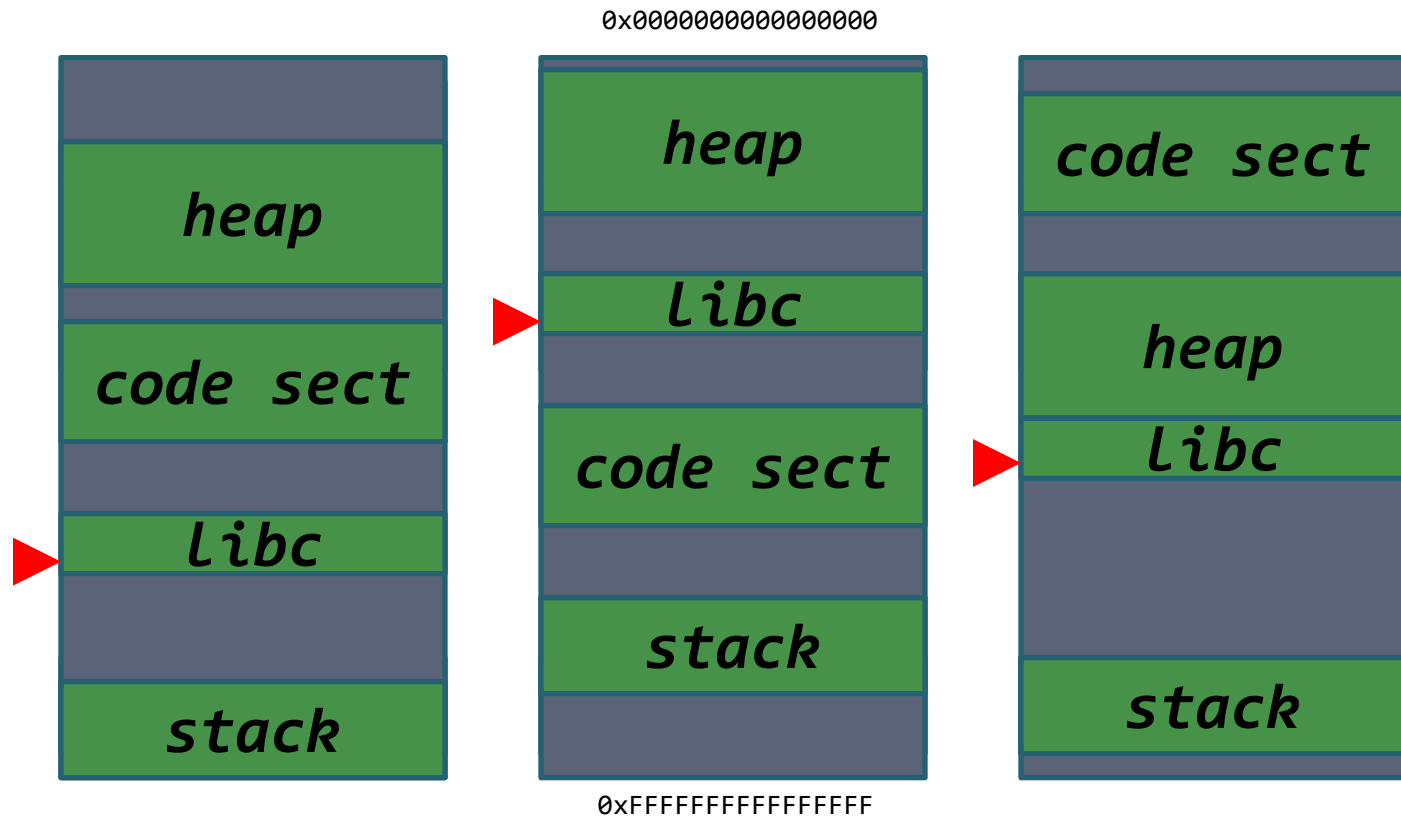
How can we defeat ASLR?

Hint:

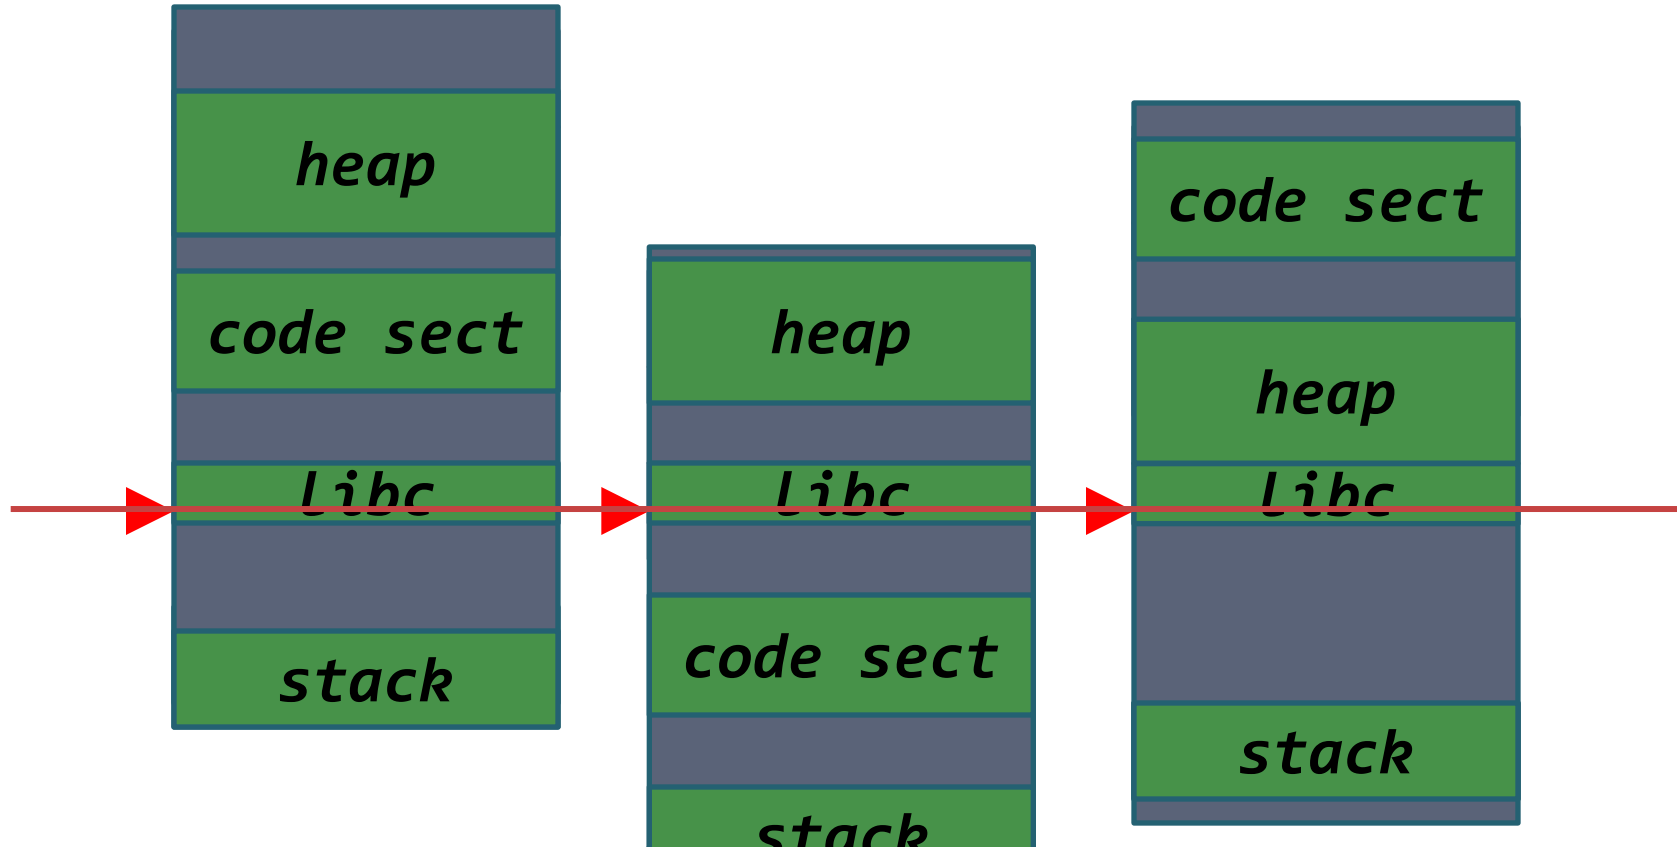
All of libc is at a single offset.

Over-read a single pointer in libc!

# Memory Layout (with ASLR)



# Memory Layout (with ASLR)



*Everything* must be relocatable to be effective

A single code section that can be referenced may provide enough ROP gadgets for exploitation

Attacker may disclose the offset of an entire chunk!

**Fine-grained ASLR** shuffles code within the chunks.



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

DEP

Extraneous function removal

ASLR

**Stack Canaries**

Defender:

Attackers keep overwriting return addresses!

Defender:

We shall protect the return address!

Keep a canary in the coal mine!

# Stack Canaries



```
# on function call:
```

```
canary = secret
```

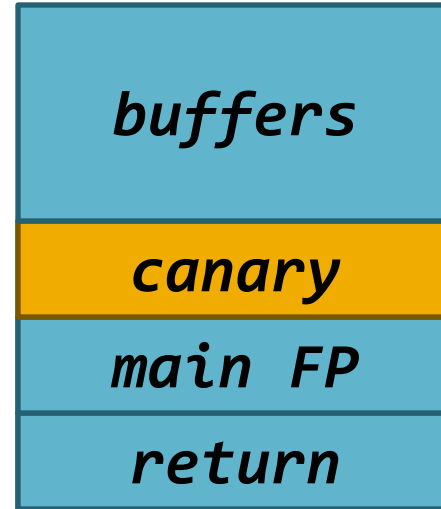
*canary*

*main FP*

*return*



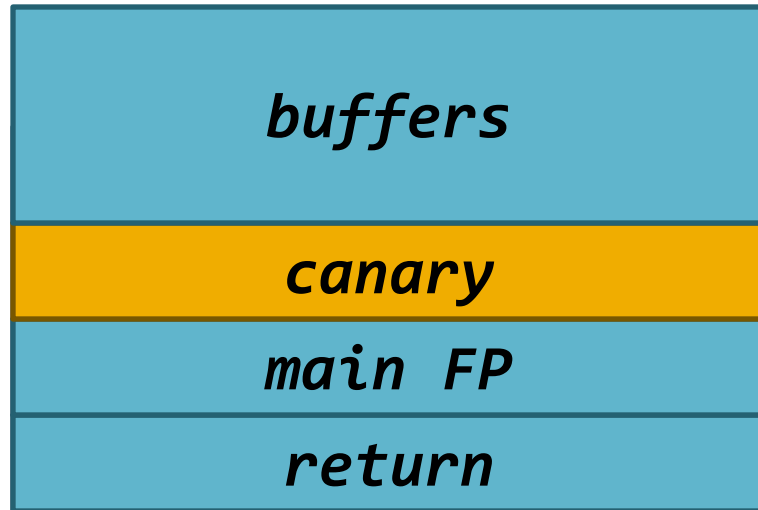
# Stack Canaries



# Stack Canaries



```
# vulnerability:  
strcpy(buffer, str)
```



```
# vulnerability:
```

```
strcpy(buffer, str)
```

**AAAAAAAA...**

**0x4141414141414141**

**0x4141414141414141**

**0x4141414141414141**

```
# on leave:
```

```
if canary != expected:  
    goto stack_chk_fail  
return
```

**AAAAAAAA...**

**0x4141414141414141**

**0x4141414141414141**

**0x4141414141414141**

# Stack Canaries



\*\*\* stack smashing detected \*\*\*

```
# on leave:
```

```
if canary != expected:  
    goto stack_chk_fail  
return
```

**AAAAAAAA...**

**0x4141414141414141**

**0x4141414141414141**

**0x4141414141414141**



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

**Buffer Over-read**

DEP

Extraneous function removal

ASLR

Stack Canaries

```
int getField(int socket, char* field){  
    int fieldLen = 0;  
    read(socket, &fieldLen, 4);  
    read(socket, field, fieldLen);  
    return fieldLen;  
}
```

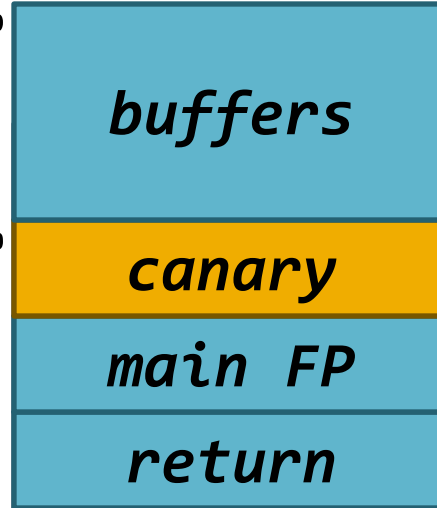

```
int sendField(int socket, char* field){  
    int fieldLen = 0;  
    read(socket, &fieldLen, 4);  
    write(socket, field, fieldLen);  
    return fieldLen;  
}
```




# Buffer Over-read



THE INTERNET



# Buffer Over-read



# Buffer Over-read



# Buffer Over-read



# Buffer Over-read



```
# on return:  
  
if canary != expected:  
    goto stack_chk_fail  
return
```



*buffers*

*canary*

*main FP*

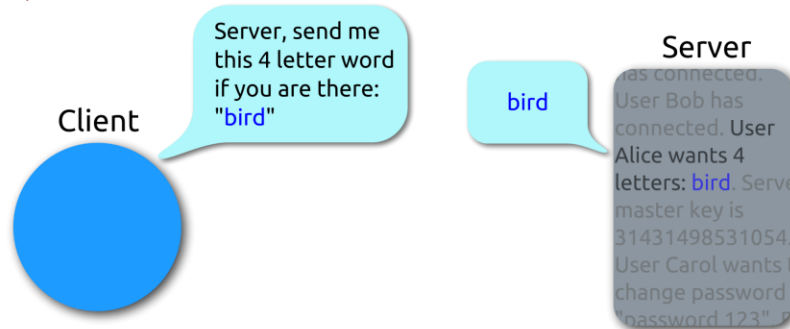
*return*

# Heartbleed

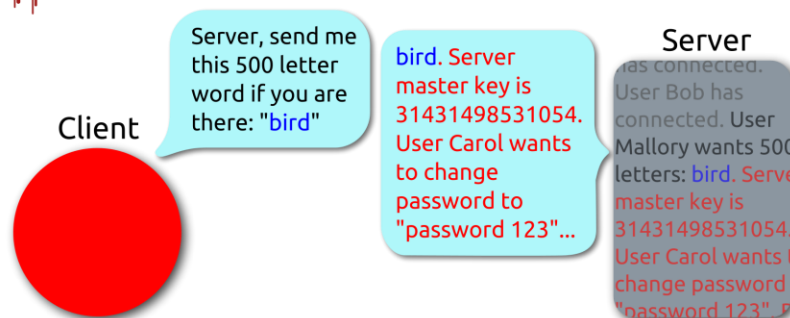




## Heartbeat – Normal usage



## Heartbeat – Malicious usage



# ...totally unrelated



**Andrew Ayer**

@\_\_agwa



This will be OpenSSL's first "CRITICAL" vulnerability since 2016. Examples of "CRITICAL" vulnerabilities include "significant disclosure of the contents of server memory (potentially revealing user details), ...



**Mark J Cox** @iamamoose · Oct 25

OpenSSL 3.0.7 update to fix Critical CVE out next Tuesday 1300-1700UTC. Does not affect versions before 3.0. [mta.openssl.org/pipermail/open...](https://mta.openssl.org/pipermail/open...)

10:37 AM · Oct 25, 2022 · Twitter Web App

318 Retweets 35 Quote Tweets 498 Likes



**Andrew Ayer** @\_\_agwa · Oct 25



Replying to @\_\_agwa

... vulnerabilities which can be easily exploited remotely to compromise server private keys or where remote code execution is considered likely in common situations" 🗑️([openssl.org/policies/gener...](https://openssl.org/policies/gener...))





# Cat-and-Mouse Exploitation

DEP

Extraneous function removal

ASLR

Stack Canaries

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

Buffer Over-read  
**Integer Casting**



# Cat-and-Mouse Exploitation

DEP

Extraneous function removal

ASLR

Stack Canaries

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

Buffer Over-read  
**Integer Overflow**

# Integer Overflow



Unsafe:

`strcpy` and friends (`str*`)

`sprintf`

`gets`

Use instead:

`strncpy` and friends (`strn*`)

`snprintf`

`fgets`

Defender:

OK fine, let's replace `strcpy` with `strncpy`

Attacker:

Find values of `n` that break `strncpy`

# Integer Overflow



```
void foo(int *array, int len) {  
    int *buf;  
    buf = malloc(len * sizeof(int));  
    if (!buf)  
        return;  
  
    int i;  
    for (i=0; i<len; i++) {  
        buf[i] = array[i];  
    }  
}
```

# Integer Overflow



```
void foo(int *array, int len) {  
    int *buf;  
    buf = malloc(len * sizeof(int));  
    if (!buf)  
        return;
```

What if len is very large?

```
    int i;  
    for (i=0; i<len; i++) {  
        buf[i] = array[i];  
    }  
}
```

# Integer Overflow



len = 1,073,742,024 ( $\approx$  1 billion)

0x400000c8

# Integer Overflow



len = 1,073,742,024 ( $\approx$  1 billion)

0x400000c8

len \* 4 = 4,294,968,096 ( $\approx$  4 billion)

0x100000320



# Integer Overflow



len = 1,073,742,024 ( $\approx$  1 billion)

0x400000c8

len \* 4 = 4,294,968,096 ( $\approx$  4 billion)

0x100000320

\*Cannot be represented in 32 bits\*

# Integer Overflow



len = 1,073,742,024 ( $\approx$  1 billion)

0x400000c8

len \* 4 = 4,294,968,096 ( $\approx$  4 billion)

0x100000320

\*Cannot be represented in 32 bits\*

0x00000320 as uint32

# Integer Overflow



`len = 1,073,742,024 (≈ 1 billion)`

`0x400000c8`

`len * 4 = 4,294,968,096 (≈ 4 billion)`

`0x100000320`

`*Cannot be represented in 32 bits*`

`0x00000320 as uint32`

`len * 4 = 800`

# Integer Overflow



```
void foo(int *array, int len) {  
    size    int *buf;  
800 → buf = malloc(len * sizeof(int));  
buffer    if (!buf)  
            return;  
  
    int i;  
    for (i=0; i<len; i++) {  
        buf[i] = array[i];  
    }  
}
```

Write  
≈ 1 billion  
elements



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

Buffer Over-read

**Signed/Unsigned Integers**

DEP

Extraneous function removal

ASLR

Stack Canaries

# Signed vs. Unsigned Integers



```
int sendField(int socket, char *field){  
    int fieldLen = 0;  
    read(socket, &fieldLen, 4);  
    write(socket, field, fieldLen);  
    return fieldLen;  
}
```

# Signed vs. Unsigned Integers



```
int sendField(int socket, char *field){
    int fieldLen = 0;
    read(socket, &fieldLen, 4);
    if (fieldLen > 10) {
        return; // Not this time :-D
    }
    write(socket, field, fieldLen);
    return fieldLen;
}
```

# Signed vs. Unsigned Integers



```
int sendField(int socket, char *field){  
    Signed type! → int fieldLen = 0;  
    read(socket, &fieldLen, 4);  
    if (fieldLen > 10) {  
        return; // Not this time :-D  
    }  
    write(socket, field, fieldLen);  
    return fieldLen;  
}
```

Negative Number



# Signed vs. Unsigned Integers



```
int sendField(int socket, char *field){  
    Signed type! → int fieldLen = 0;  
    read(socket, &fieldLen, 4);  
    if (fieldLen > 10) {  
        return; // Not this time :-D  
    }  
    write(socket, field, fieldLen);  
    return fieldLen;  
}
```

Negative Number

Passes Signed Check

# Signed vs. Unsigned Integers



```
int sendField(int socket, char *field){  
    Signed type! → int fieldLen = 0;  
    read(socket, &fieldLen, 4);  
    if (fieldLen > 10) {  
        return; // Not this time :-D  
    }  
    write(socket, field, fieldLen);  
    return fieldLen;  
}
```

Negative Number

Passes Signed Check

Treated as a very large number  
(third argument is type size\_t, an unsigned integer)

# Signed vs. Unsigned Integers



fieldLen as size\_t

fieldLen as int

Enter binary number

10101101101011011010110110101101 2

= Convert × Reset ↕ Swap

Decimal number (10 digits)

2913840557 10

Decimal from signed 2's complement (10 digits)

-1381126739 10

The SAME bits give us either a HUGE number (unsigned) or a TINY number (signed)  
It depends on how the bits are *interpreted*



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

Buffer Over-read  
Integer Casting

DEP

Extraneous function removal

ASLR

Stack Canaries

**Automated Testing**



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

Buffer Over-read  
Integer Casting  
Automated Testing

DEP

Extraneous function removal

ASLR

Stack Canaries

**Automated Testing**

# Automated Testing



The Problem With Computers:

Vulnerabilities are hard to find by hand

Defender:

(and attacks use them ☹️)

Attacker:

(and attacks use them 😊)

Defender and Attacker (*in unison*):

Automate the process!

# Automated Testing



Finding vulnerabilities manually is very hard

If source is available:

- Tons of potential vulnerabilities in code base

If closed source:

- Reverse Engineering is laborious

## Memory Analysis Tools

Incredibly useful for finding memory leaks

Execute in a virtual environment

& perform dynamic run-time checks

Does the program access uninitialized memory?

Does the program use memory after it's freed?



## Static Analysis Tools

- Look for dangerous coding patterns and practices

- Usually requires complete source code

- Large number of false-positives

- Are integers mixing signed and unsigned usage?

- Are all variables initialized when declared?

## Taint Analysis Tools

- Trace value usage throughout code

- Attempt to identify when untrusted data is used

- Is a user-supplied value used to index an array?

- Is an unsafe value used to shell-out?

## Fuzzers

“Brute Force Testing”

Generate inputs and monitor program’s behavior

More advanced optimize for code coverage

If I give you really long strings, will you crash?

If I give you random data, will you crash?

If I give you broken formats, will you crash?



# Cat-and-Mouse Exploitation

Buffer Overflow  
Stack Shellcode

Return-to-libc

ROP

Buffer Over-read  
Integer Casting

DEP

Extraneous function removal

ASLR

Stack Canaries

Automated Testing

**Toolbox of Exploitation Techniques**

Every vulnerability is different

Some are not exploitable at all

Sometimes it takes multiple bugs to create an exploit (“Bug Chains”)

Buffer over-read (canary) + Buffer over-read (ASLR reference) + Buffer overflow (load exploit) + ROP chain (disable DEP) + Jump to shellcode

# Taking the Easy Road



Don't overly complicate the exploit

Is there an n-day?

Can you exploit a function without canaries?

Can you pivot from another application?

Can you brute-force a canary?

# Use After Free



Common in multi-threaded programs that share variables


Though can exist in single threaded programs

Sometimes caused by a race condition


# Use After Free



Heap:




Stack:







# Use After Free




# Use After Free




# Use After Free



# Use After Free



# Use After Free



# Format String Vulnerability



Attack programmer's lack of sanitization

```
printf("%s\n", argv[1]);
```

```
printf(argv[1]);
```

← Number of arguments  
determined by the  
format of the string

```
./a.out "%p %p %p %p %p %p"
```

# Further Reading



- “Smashing the Stack 21<sup>st</sup> Century”  
<https://thesquareplanet.com/blog/smashing-the-stack-21st-century/>
- Blexim’s “Basic Integer Overflows”  
<http://www.phrack.org/issues.html?issue=60&id=10>
- irOnstone’s “Binary Exploitation Notes”  
<https://irOnstone.gitbook.io/notes/>
- Aleph One’s “Smashing the Stack for Fun and Profit” (uses 32-bit instead of x64, still a classic!) <http://insecure.org/stf/smashstack.html>

Note: further reading may not use x86-64 with Intel syntax, so be mindful of differences

# Coming Up



**Project 4 due November 14 at 6 p.m.**

**Lab 4 due October 31 at 6 p.m. (next Thursday)**



**Tuesday, Oct. 29**

**Malware**



**Thursday, Oct. 31**

**Access Control and  
Isolation**

